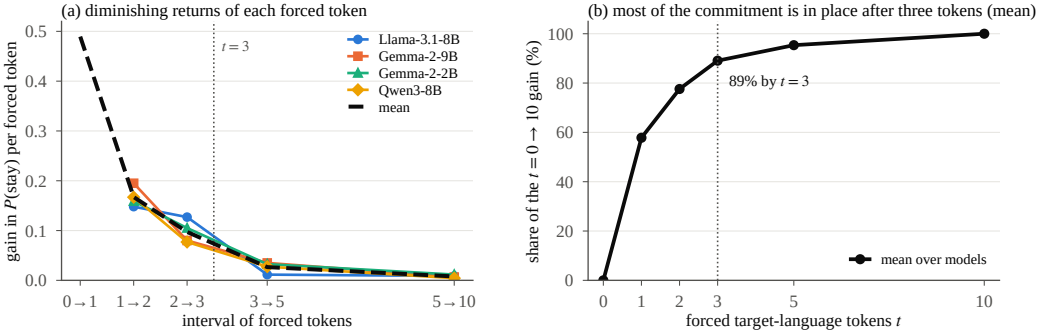
(a) diminishing returns of each forced token
gain in P(stay) per forced token
0.5
0.4
0.3
0.2
0.1
0.0
t = 3
Llama-3.1-8B
Gemma-2-9B
Gemma-2-2B
Qwen3-8B
mean
0→1
1→2
2→3
3→5
5→10
interval of forced tokens
(b) most of the commitment is in place after three tokens (mean)
share of the t = 0 → 10 gain (%)
100
80
60
40
20
0
89% by t = 3
mean over models
0
1
2
3
5
10
forced target-language tokens t